

| | |
|-----------------|-----------------|
| Assignment Date | 01-October-2022 |
| Student Name | MONIKA K |
| Student Roll No | 61771931034 |
| Maximum marks | 2 marks |

Assignment -3 Python Programming

1. Download the dataset

```
In [1]: #importing the Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

2. Load the dataset into the tool.

```
In [2]: #Loading the dataset

d = pd.read_csv(r'Downloads/abalone.csv')
```

3. Perform Below Visualizations.

• Univariate Analysis

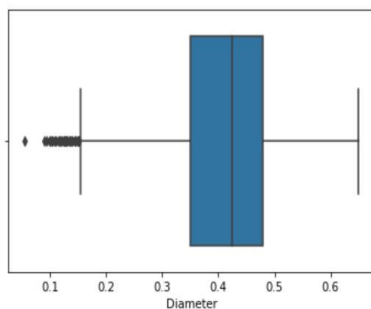
```
In [3]: d.head()
```

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 15 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 7 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 9 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 10 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 7 |

```
In [4]: #Boxplot

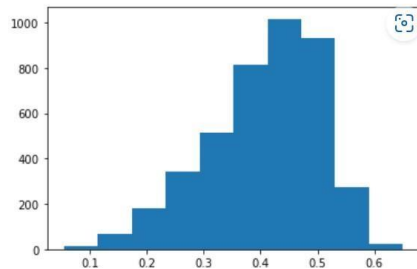
sns.boxplot(d['Diameter'])

Out[4]: <AxesSubplot:xlabel='Diameter'>
```



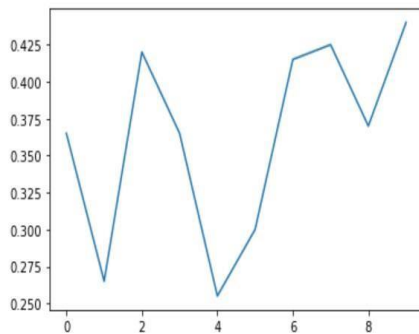
```
In [5]: #histogram
plt.hist(d['Diameter'])
```

```
Out[5]: (array([ 13.,  66., 180., 344., 513., 812., 1017., 934., 275.,
                23.]),
array([0.055, 0.1145, 0.174, 0.2335, 0.293, 0.3525, 0.412, 0.4715,
        0.531, 0.5905, 0.65 ]),
<BarContainer object of 10 artists>)
```



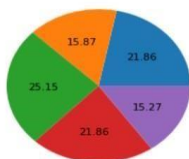
```
In [6]: #Line plot
plt.plot(d['Diameter'].head(10))
```

```
Out[6]: [<matplotlib.lines.Line2D at 0x1c2ed71d130>]
```



```
In [7]: #piechart
plt.pie(d['Diameter'].head(), autopct='%2f')
```

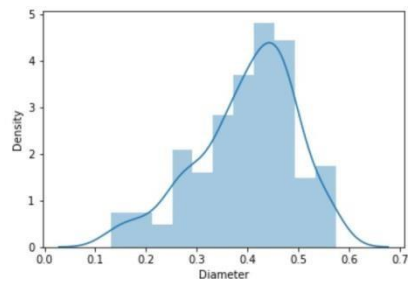
```
Out[7]: ([<matplotlib.patches.Wedge at 0x1c2ed77fd00>,
<matplotlib.patches.Wedge at 0x1c2ed78d460>,
<matplotlib.patches.Wedge at 0x1c2ed78db80>,
<matplotlib.patches.Wedge at 0x1c2ed7992e0>,
<matplotlib.patches.Wedge at 0x1c2ed799a30>],
[Text(0.8507215626110558, 0.6973326486753676, ''),
Text(-0.32611344931648134, 1.0505474849691026, ''),
Text(-1.0998053664078908, -0.02069193128747144, ''),
Text(-0.08269436219656089, -1.096887251480709, ''),
Text(0.9758446362287218, -0.5076684409569241, ''),
Text(0.464029943242394, 0.3803632629138369, '21.86'),
Text(-0.17788006326353525, 0.5730259008922377, '15.87'),
Text(-0.5998938362224858, -0.011286507974984419, '25.15'),
Text(-0.045106015743578656, -0.5983021371712958, '21.86'),
Text(0.5322788924883937, -0.2769100587037768, '15.27')])
```



```
In [8]: #distplot
```

```
sns.distplot(d['Diameter'].head(200))
```

```
Out[8]: <AxesSubplot:xlabel='Diameter', ylabel='Density'>
```




● Bi - Variate Analysis

```
In [9]: #scatter plot
```

```
plt.scatter(d['Diameter'].head(500), d['Length'].head(500))
```

```
Out[9]: <matplotlib.collections.PathCollection at 0x1c2edcc2d60>
```



```
In [10]: #bar plot
```

```
plt.bar(d['Sex'].head(10), d['Rings'].head(10))
```


```
#Labelling of x,y and result
```

```
plt.title('Bar plot')
```

```
plt.xlabel('Diameter')
```


```
plt.ylabel('Rings')
```

```
Out[10]: Text(0, 0.5, 'Rings')
```




```
In [11]: sns.barplot(d['Sex'], d['Rings'])
```

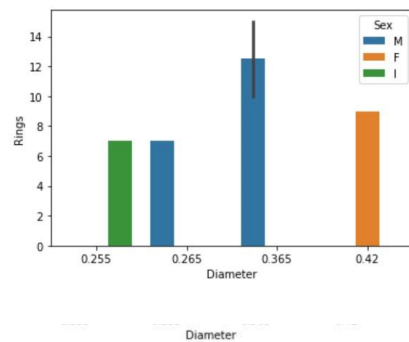
```
Out[11]: <AxesSubplot:xlabel='Sex', ylabel='Rings'>
```



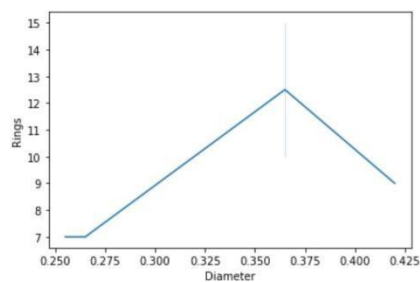
```
In [12]: #joint plot
sns.jointplot(d['Diameter'].head(50),d['Rings'].head(50))
Out[12]: <seaborn.axisgrid.JointGrid at 0x1c2edde3160>
```



```
In [13]: #bar plot
sns.barplot('Diameter','Rings',hue='Sex',data=d.head())
Out[13]: <AxesSubplot:xlabel='Diameter', ylabel='Rings'>
```



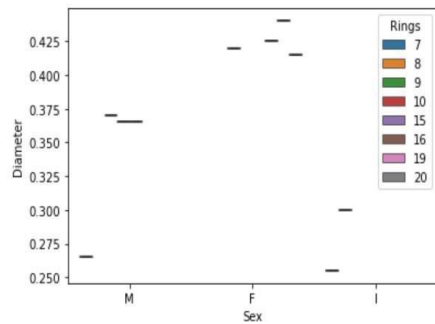
```
In [14]: sns.lineplot(d['Diameter'].head(),d['Rings'].head())
Out[14]: <AxesSubplot:xlabel='Diameter', ylabel='Rings'>
```



• Multi - Variate Analysis


```
In [15]: #boxplot
sns.boxplot(d['Sex'],d['Diameter'].head(10),d['Rings'].head(10))
```

```
Out[15]: <AxesSubplot:xlabel='Sex', ylabel='Diameter'>
```



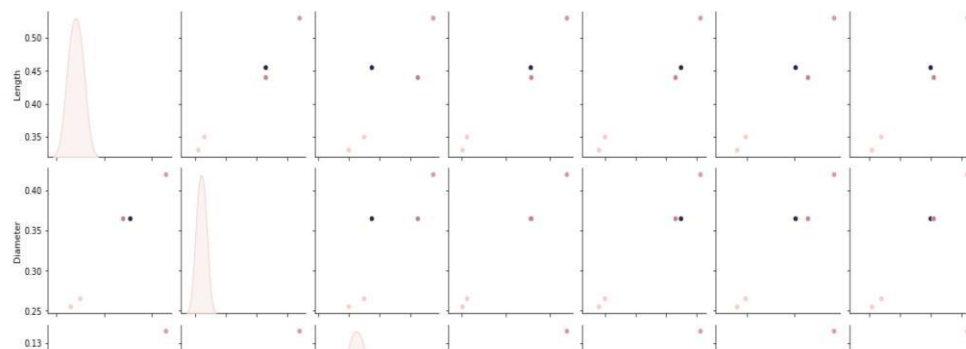
```
In [16]: #heat map
fig=plt.figure(figsize=(8,5))
sns.heatmap(d.head().corr(),annot=True)
```

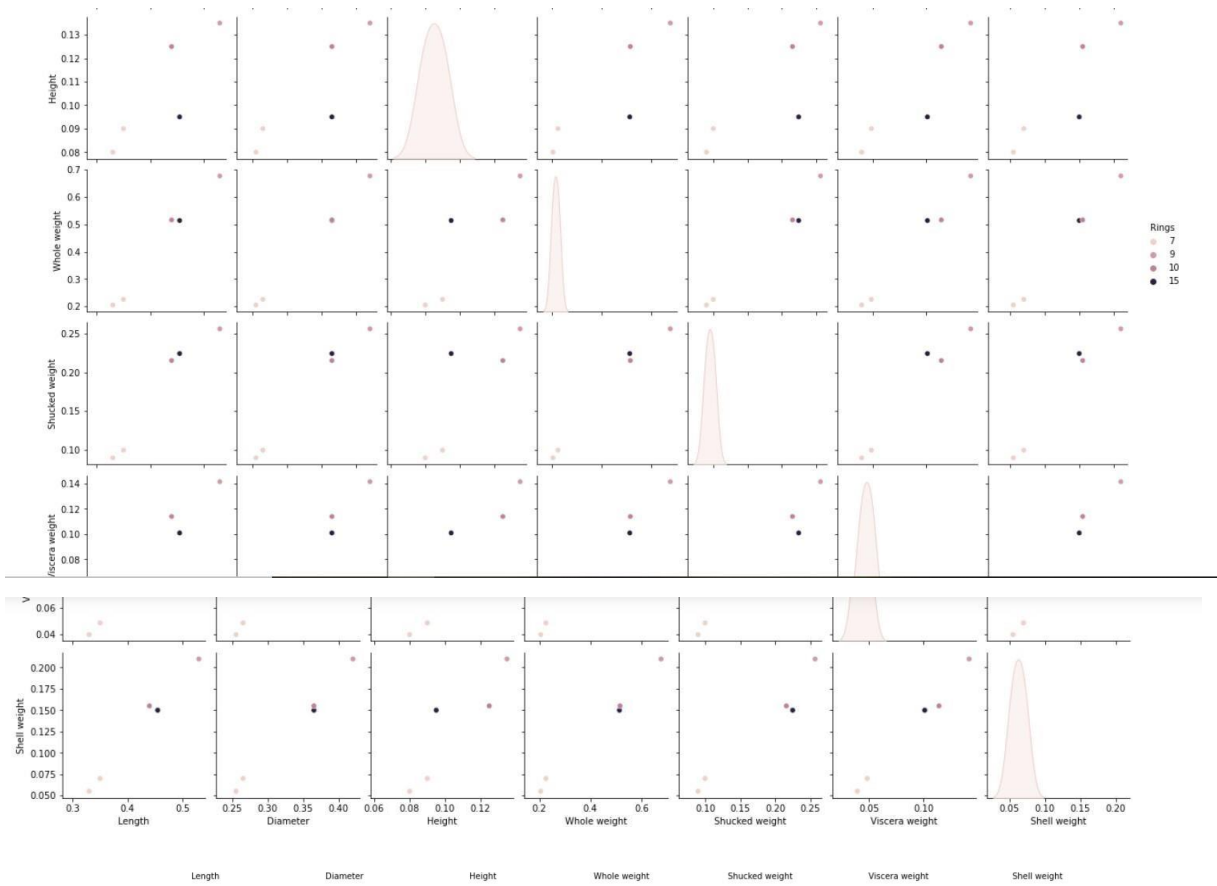
```
Out[16]: <AxesSubplot:~>
```



```
In [17]: #pair plot
sns.pairplot(d.head(),hue='Rings')
```


```
Out[17]: <seaborn.axisgrid.PairGrid at 0x1c2edd07fd0>
```

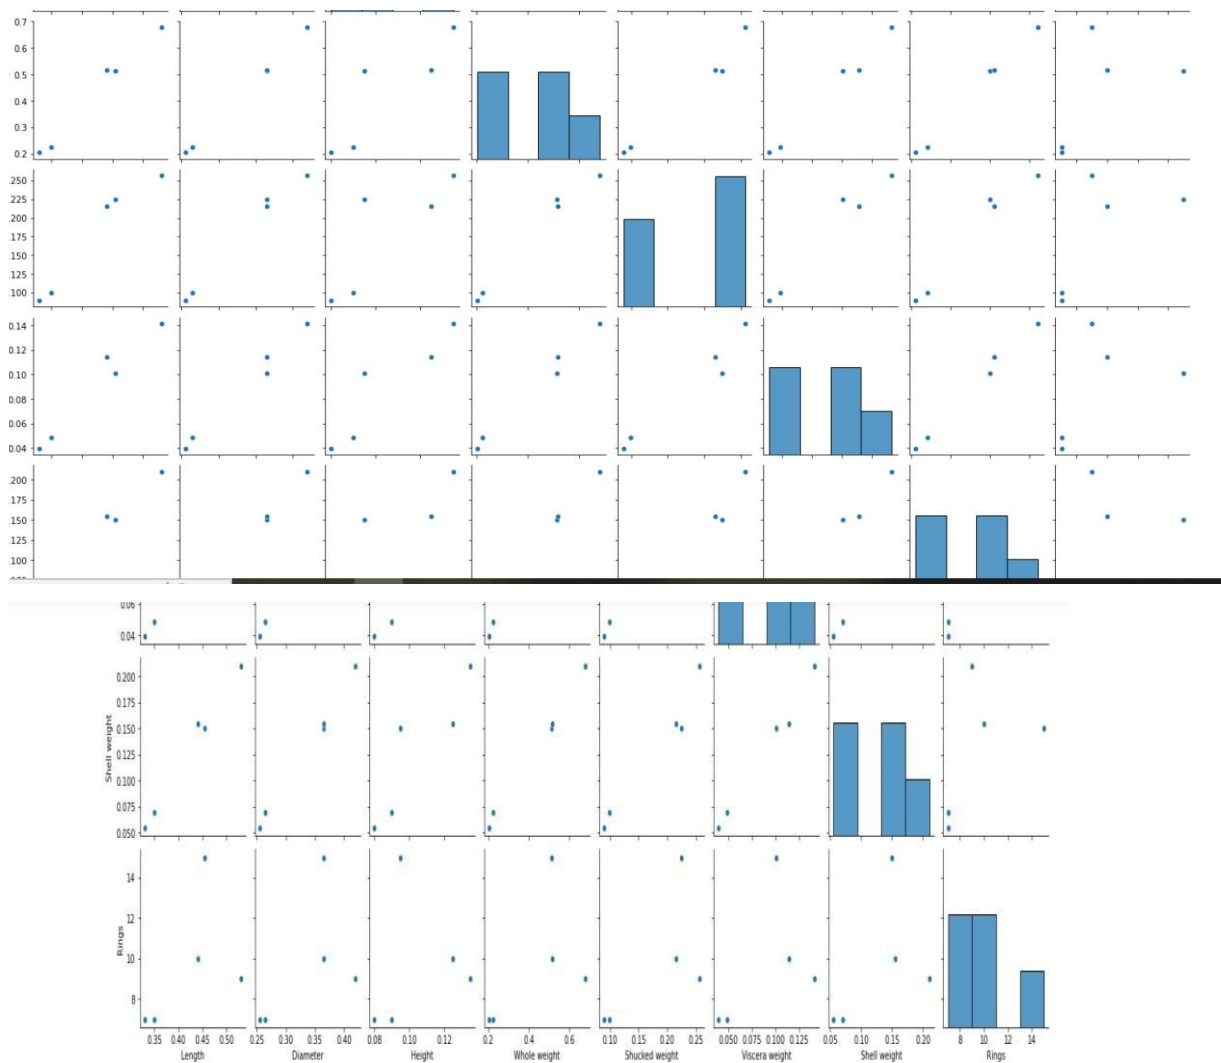




```
In [18]: sns.pairplot(d.head())
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x1c2f243e910>
```





4. Perform descriptive statistics on the dataset.

In [19]: `#head`

`d.head()`

Out[19]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 15 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 7 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 9 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 10 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 7 |

In [20]: `#tail`

`d.tail()`

Out[20]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 4172 | F | 0.565 | 0.450 | 0.165 | 0.8870 | 0.3700 | 0.2390 | 0.2490 | 11 |
| 4173 | M | 0.590 | 0.440 | 0.135 | 0.9660 | 0.4390 | 0.2145 | 0.2605 | 10 |

| | | | | | | | | | |
|------|---|-------|-------|-------|--------|--------|--------|--------|----|
| 4174 | M | 0.600 | 0.475 | 0.205 | 1.1760 | 0.5255 | 0.2875 | 0.3080 | 9 |
| 4175 | F | 0.625 | 0.485 | 0.150 | 1.0945 | 0.5310 | 0.2610 | 0.2960 | 10 |
| 4176 | M | 0.710 | 0.555 | 0.195 | 1.9485 | 0.9455 | 0.3765 | 0.4950 | 12 |

In [21]: d.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight    4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Rings            4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

In [22]: d.describe()

Out[22]:

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|-------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 |
| mean | 0.523992 | 0.407881 | 0.139516 | 0.828742 | 0.359367 | 0.180594 | 0.238831 | 9.933684 |
| std | 0.120093 | 0.099240 | 0.041827 | 0.490389 | 0.221963 | 0.109614 | 0.139203 | 3.224169 |
| min | 0.075000 | 0.055000 | 0.000000 | 0.002000 | 0.001000 | 0.000500 | 0.001500 | 1.000000 |
| 25% | 0.450000 | 0.350000 | 0.115000 | 0.441500 | 0.186000 | 0.093500 | 0.130000 | 8.000000 |
| 50% | 0.545000 | 0.425000 | 0.140000 | 0.799500 | 0.336000 | 0.171000 | 0.234000 | 9.000000 |
| 75% | 0.615000 | 0.480000 | 0.165000 | 1.153000 | 0.502000 | 0.253000 | 0.329000 | 11.000000 |
| max | 0.815000 | 0.650000 | 1.130000 | 2.825500 | 1.488000 | 0.760000 | 1.005000 | 29.000000 |

In [23]: #mode

d.mode().T

Out[23]:

| | 0 | 1 |
|----------------|--------|-------|
| Sex | M | NaN |
| Length | 0.55 | 0.625 |
| Diameter | 0.45 | NaN |
| Height | 0.15 | NaN |
| Whole weight | 0.2225 | NaN |
| Shucked weight | 0.175 | NaN |
| Viscera weight | 0.1715 | NaN |
| Shell weight | 0.275 | NaN |
| Rings | 9.0 | NaN |

In [24]: d.shape

Out[24]: (4177, 9)


```
In [25]: #kurtosis
d.kurt()
```

```
Out[25]: Length      0.064621
Diameter    -0.045476
Height      76.025509
Whole weight -0.023644
Shucked weight 0.595124
Viscera weight 0.084012
Shell weight 0.531926
Rings       2.330687
dtype: float64
```

```
In [26]: #skewness
d.skew()
```

```
Out[26]: Length      -0.639873
Diameter    -0.609198
Height       3.128817
Whole weight 0.530959
Shucked weight 0.719098
Viscera weight 0.591852
Shell weight 0.620927
Rings       1.114102
dtype: float64
```

```
In [27]: #variance
d.var()
```

```
Out[27]: Length      0.014422
Diameter    0.009849
Height      0.001750
Whole weight 0.240481
Shucked weight 0.049268
Viscera weight 0.012015
Shell weight 0.019377
Rings      10.395266
dtype: float64
```

```
In [28]: #finding unique values for columns
d.nunique()
```

```
Out[28]: Sex          3
Length        134
Diameter       111
Height         51
Whole weight  2429
Shucked weight 1515
Viscera weight 880
Shell weight  926
Rings         28
dtype: int64
```

5. Check for Missing values and deal with them.

```
In [29]: #finding missing values
d.isna()
```

```
Out[29]:
```

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|-------|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | False | False | False | False | False | False | False | False | False |
| 1 | False | False | False | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4172 | False | False | False | False | False | False | False | False | False |
| 4173 | False | False | False | False | False | False | False | False | False |
| 4174 | False | False | False | False | False | False | False | False | False |
| 4175 | False | False | False | False | False | False | False | False | False |
| 4176 | False | False | False | False | False | False | False | False | False |

4177 rows x 9 columns

```
In [30]: d.isna().any()
```

```
Out[30]: Sex           False
Length           False
Diameter         False
Height           False
Whole weight     False
Shucked weight   False
Viscera weight   False
Shell weight     False
Rings            False
dtype: bool
```

```
In [31]: d.isna().sum()
```

```
Out[31]: Sex           0
Length           0
Diameter         0
Height           0
Whole weight     0
Shucked weight   0
Viscera weight   0
Shell weight     0
Rings            0
dtype: int64
```

```
dtype: bool
```

```
In [31]: d.isna().sum()
```

```
Out[31]: Sex           0
Length           0
Diameter         0
Height           0
Whole weight     0
Shucked weight   0
Viscera weight   0
Shell weight     0
Rings            0
dtype: int64
```

```
In [32]: d.isna().any().sum()
#no missing values
```


```
Out[32]: 0
```

6. Find the outliers and replace them outliers

```
In [33]: #finding outliers
```

```
sns.boxplot(d['Diameter'])
```

```
Out[33]: <AxesSubplot: xlabel='Diameter'>
```



```
In [34]: #handling outliers
```

```
qnt=d.quantile(q=[0.25,0.75])
qnt
```

```
Out[34]:
```

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0.25 | 0.450 | 0.35 | 0.115 | 0.4415 | 0.186 | 0.0935 | 0.130 | 8.0 |
| 0.75 | 0.615 | 0.48 | 0.165 | 1.1530 | 0.502 | 0.2530 | 0.329 | 11.0 |

```
In [35]: iqr=qnt.loc[0.75]-qnt.loc[0.25]
```

```
iqr
```

```
Out[35]: Length           0.1650
Diameter           0.1300
Height           0.0500
Whole weight       0.7115
Shucked weight     0.3160
Viscera weight     0.1595
Shell weight       0.1990
Rings              3.0000
dtype: float64
```

```
In [36]: lower=qnt.loc[0.25]-(1.5*iqr)
lower
```

```
Out[36]: Length      0.20250
Diameter    0.15500
Height      0.04000
Whole weight -0.62575
Shucked weight -0.28800
Viscera weight -0.14575
Shell weight -0.16850
Rings       3.50000
dtype: float64
```

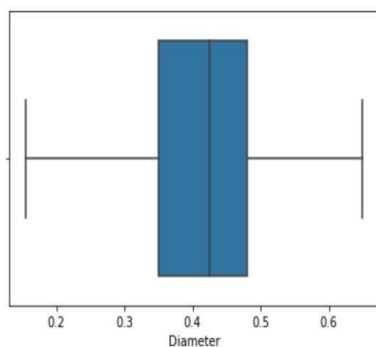
```
In [37]: upper=qnt.loc[0.75]+(1.5*iqr)
upper
```

```
Out[37]: Length      0.86250
Diameter    0.67500
Height      0.24000
Whole weight 2.22025
Shucked weight 0.97600
Viscera weight 0.49225
Shell weight 0.62750
Rings      15.50000
dtype: float64
```

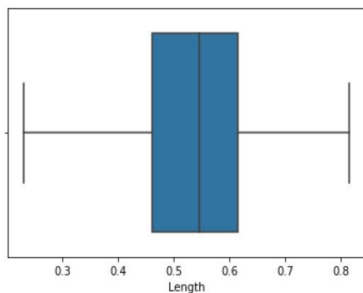
```
In [38]: # replacing outliers

##Diameter
d['Diameter']=np.where(d['Diameter']<0.155,0.4078,d['Diameter'])
sns.boxplot(d['Diameter'])
```

```
Out[38]: <AxesSubplot:xlabel='Diameter'>
```



```
Out[41]: <AxesSubplot:xlabel='Length'>
```



```
In [42]: ## Height
sns.boxplot(d['Height'])
```

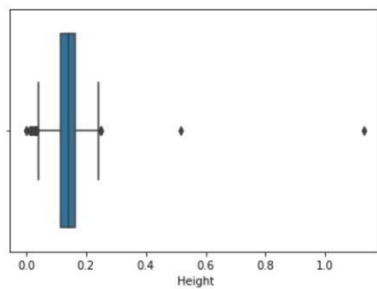
```
Out[42]: <AxesSubplot:xlabel='Height'>
```



```
In [42]: ## Height
```

```
sns.boxplot(d['Height'])
```

```
Out[42]: <AxesSubplot:xlabel='Height'>
```

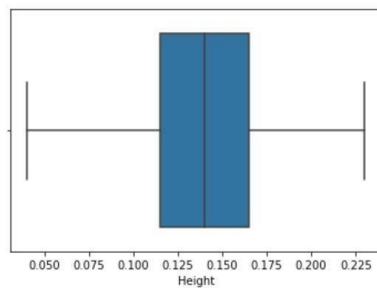


```
In [43]: d['Height']=np.where(d['Height']<0.04,0.139, d['Height'])  
d['Height']=np.where(d['Height']>0.23,0.139, d['Height'])
```

```
d['Height']=np.where(d['Height']>0.23,0.139, d['Height'])
```

```
In [44]: sns.boxplot(d['Height'])
```

```
Out[44]: <AxesSubplot:xlabel='Height'>
```

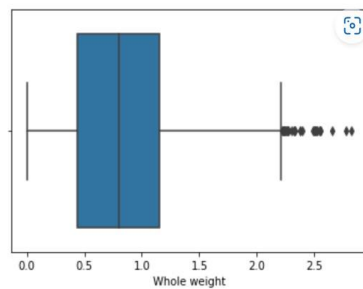


```
In [45]: ## Whole weight
```

```
sns.boxplot(d['Whole weight'])
```

```
Out[45]: <AxesSubplot:xlabel='Whole weight'>
```

```
Out[45]: <AxesSubplot:xlabel='Whole weight'>
```



```
In [46]: d['Whole weight']=np.where(d['Whole weight']>0.9,0.82, d['Whole weight'])
```


```
In [47]: sns.boxplot(d['Whole weight'])
```

```
Out[47]: <AxesSubplot:xlabel='Whole weight'>
```



```
In [47]: sns.boxplot(d['Whole weight'])
```

```
Out[47]: <AxesSubplot:xlabel='Whole weight'>
```



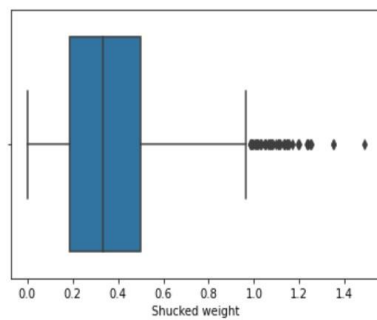
```
In [48]: ## Shucked weight
```

```
sns.boxplot(d['Shucked weight'])
```

```
Out[48]: <AxesSubplot:xlabel='Shucked weight'>
```



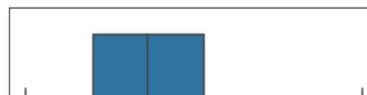
```
Out[48]: <AxesSubplot:xlabel='Shucked weight'>
```



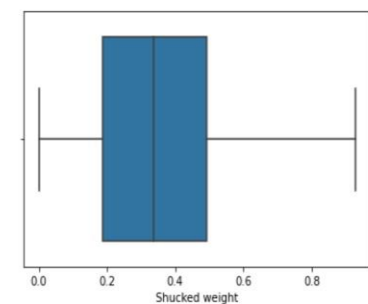
```
In [49]: d['Shucked weight']=np.where(d['Shucked weight']>0.93,0.35, d['Shucked weight'])
```

```
In [50]: sns.boxplot(d['Shucked weight'])
```

```
Out[50]: <AxesSubplot:xlabel='Shucked weight'>
```



```
Out[50]: <AxesSubplot:xlabel='Shucked weight'>
```



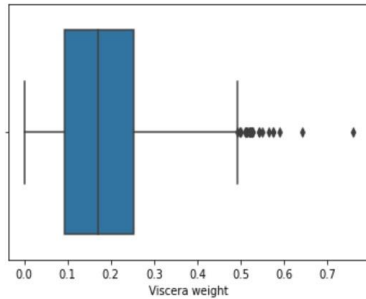
```
In [51]: ## Viscera weight
```

```
sns.boxplot(d['Viscera weight'])
```

```
Out[51]: <AxesSubplot:xlabel='Viscera weight'>
```



```
Out[51]: <AxesSubplot:xlabel='Viscera weight'>
```



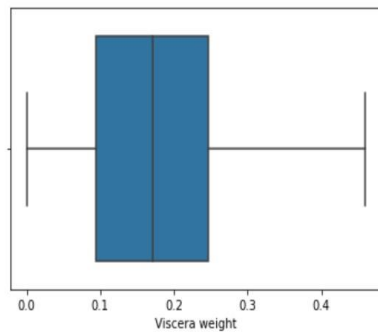
```
In [52]: d['Viscera weight']=np.where(d['Viscera weight']>0.46,0.18, d['Viscera weight'])
```

```
In [53]: sns.boxplot(d['Viscera weight'])
```

```
Out[53]: <AxesSubplot:xlabel='Viscera weight'>
```

```
In [53]: sns.boxplot(d['Viscera weight'])
```

```
Out[53]: <AxesSubplot:xlabel='Viscera weight'>
```



```
In [54]: ## Shell weight
```

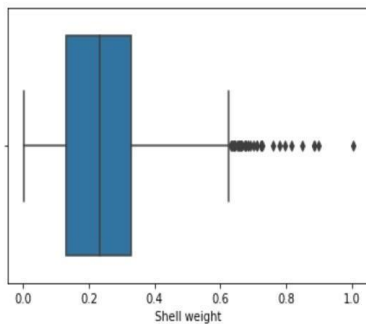
```
sns.boxplot(d['Shell weight'])
```

```
Out[54]: <AxesSubplot:xlabel='Shell weight'>
```

```
In [54]: ## Shell weight
```

```
sns.boxplot(d['Shell weight'])
```

```
Out[54]: <AxesSubplot:xlabel='Shell weight'>
```




```
In [55]: d['Shell weight']=np.where(d['Shell weight']>0.61,0.2388, d['Shell weight'])
```

```
In [56]: sns.boxplot(d['Shell weight'])
```

```
Out[56]: <AxesSubplot:xlabel='Shell weight'>
```

```
In [56]: sns.boxplot(d['Shell weight'])
```

```
Out[56]: <AxesSubplot:xlabel='Shell weight'>
```



7. Check for Categorical columns and perform encoding.

```
In [57]: #one hot encoding
```

```
d['Sex'].replace({'M':1,'F':0,'I':2},inplace=True)
d
```

```
Out[57]:
```

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | 1 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.1500 | 15 |
| 1 | 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.0700 | 7 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.2100 | 9 |
| 3 | 1 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.1550 | 10 |
| 4 | 2 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.0550 | 7 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4172 | 0 | 0.565 | 0.450 | 0.165 | 0.8870 | 0.3700 | 0.2390 | 0.2490 | 11 |
| 4173 | 1 | 0.590 | 0.440 | 0.135 | 0.8200 | 0.4390 | 0.2145 | 0.2605 | 10 |
| 4174 | 1 | 0.600 | 0.475 | 0.205 | 0.8200 | 0.5255 | 0.2875 | 0.3080 | 9 |
| 4175 | 0 | 0.625 | 0.485 | 0.150 | 0.8200 | 0.5310 | 0.2610 | 0.2960 | 10 |
| 4176 | 1 | 0.710 | 0.555 | 0.195 | 0.8200 | 0.3500 | 0.3765 | 0.4950 | 12 |

4177 rows x 9 columns

8. Split the data into dependent and independent variables.

```
In [58]: x=d.drop(columns= ['Rings'])
y=d['Rings']
x
```

```
Out[58]:
```

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|
| 0 | 1 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.1500 |
| 1 | 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.0700 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.2100 |
| 3 | 1 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.1550 |
| 4 | 2 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.0550 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4172 | 0 | 0.565 | 0.450 | 0.165 | 0.8870 | 0.3700 | 0.2390 | 0.2490 |
| 4173 | 1 | 0.590 | 0.440 | 0.135 | 0.8200 | 0.4390 | 0.2145 | 0.2605 |
| 4174 | 1 | 0.600 | 0.475 | 0.205 | 0.8200 | 0.5255 | 0.2875 | 0.3080 |
| 4175 | 0 | 0.625 | 0.485 | 0.150 | 0.8200 | 0.5310 | 0.2610 | 0.2960 |
| 4176 | 1 | 0.710 | 0.555 | 0.195 | 0.8200 | 0.3500 | 0.3765 | 0.4950 |

4177 rows x 8 columns

4177 rows x 8 columns

In [59]: y

```
Out[59]: 0      15
         1       7
         2       9
         3      10
         4       7
         ..
        4172    11
        4173    10
        4174     9
        4175    10
        4176    12
        Name: Rings, Length: 4177, dtype: int64
```

9. Scale the independent variables

In [60]: `from sklearn.preprocessing import scale #StandardScaler`

In [61]: `#Scaling the independent variables`

```
x = scale(x)
x
```

```
Out[61]: array([[ -0.0105225, -0.67088921, -0.50179694, ..., -0.61037964,
        -0.7328165, -0.64358742],
        [ -0.0105225, -1.61376082, -1.57304487, ..., -1.22513334,
        -1.24343929, -1.25742181],
        [ -1.26630752,  0.00259051,  0.08738942, ..., -0.45300269,
        -0.33890749, -0.18321163],
        ...,
        [ -0.0105225,  0.63117159,  0.67657577, ...,  0.86994729,
         1.08111018,  0.56873549],
        [ -1.26630752,  0.85566483,  0.78370057, ...,  0.89699645,
         0.82336724,  0.47666033],
        [ -0.0105225,  1.61894185,  1.53357412, ...,  0.00683308,
         1.94673739,  2.00357336]])
```

10. Split the data into training and testing

In [62]: `from sklearn.model_selection import train_test_split`

In [63]: `#splitting data to train and test`

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2)
print(x_train.shape, x_test.shape)

(3341, 8) (836, 8)
```

11. Build the Model

In [64]: `#Multiple Regression`

```
from sklearn.linear_model import LinearRegression

MLR=LinearRegression()
```

12. Train the Model

12. Train the Model

```
In [65]: MLR.fit(x_train,y_train)
```

```
Out[65]: LinearRegression()
```

13. Test the Model

```
In [66]: #prediction on the test data  
y_pred=MLR.predict(x_test)  
y_pred
```

```
Out[66]: array([11.46655124,  9.21660091,  6.59967857,  7.81824648, 12.18984569,  
11.44220895, 11.20545145,  8.71621092, 10.98237601,  6.83381457,  
10.46227495,  9.10809044, 12.39359143, 14.54491772, 13.54791716,  
10.12045364, 11.48597397,  7.73511543, 12.86466796,  8.37939955,  
 6.51920876,  8.16682072,  8.05416099, 10.22713858, 10.57995698,  
11.31009826,  7.52742935,  9.88582514, 11.25644638, 11.38973324,  
10.95569239, 10.28552912, 10.40475249, 10.44887526, 11.03343746,  
10.15213587, 10.04733695,  6.54448931, 11.86305246,  6.73817965,  
 4.07354447, 11.09033543,  7.69897797,  9.56311429, 11.63006462,  
13.17063754,  6.34451832,  7.27896893, 15.31511539,  6.92860099,  
 3.63485054,  6.80184256, 11.451762  , 10.69664795,  8.59383781,  
 7.50446583, 10.33994154, 11.85072027, 13.544946  , 10.27236403,
```

13. Test the Model

```
In [66]: #prediction on the test data  
y_pred=MLR.predict(x_test)  
y_pred
```

```
Out[66]: array([11.46655124,  9.21660091,  6.59967857,  7.81824648, 12.18984569,  
11.44220895, 11.20545145,  8.71621092, 10.98237601,  6.83381457,  
10.46227495,  9.10809044, 12.39359143, 14.54491772, 13.54791716,  
10.12045364, 11.48597397,  7.73511543, 12.86466796,  8.37939955,  
 6.51920876,  8.16682072,  8.05416099, 10.22713858, 10.57995698,  
11.31009826,  7.52742935,  9.88582514, 11.25644638, 11.38973324,  
10.95569239, 10.28552912, 10.40475249, 10.44887526, 11.03343746,  
10.15213587, 10.04733695,  6.54448931, 11.86305246,  6.73817965,  
 4.07354447, 11.09033543,  7.69897797,  9.56311429, 11.63006462,  
13.17063754,  6.34451832,  7.27896893, 15.31511539,  6.92860099,  
 3.63485054,  6.80184256, 11.451762  , 10.69664795,  8.59383781,  
 7.50446583, 10.33994154, 11.85072027, 13.544946  , 10.27236403,  
 9.18410191,  7.7208794  , 12.33421272,  6.527156  , 11.17483778,  
 7.97617745,  9.31452692,  9.56473016,  9.51077399, 12.20917888,  
12.10672271,  4.70427674,  6.38943267, 10.02410014, 11.97786002,  
12.77246335,  6.50139525, 10.64829499,  7.7058727  ,  6.05475715,  
11.28248424, 10.75341994, 17.22835762,  9.53819376,  8.96368426,  
 6.61412036, 12.00162611,  5.85400348,  4.07058709, 10.08426584,  
10.15760235, 11.50892785, 10.58412873, 10.32113545, 12.98841501,
```

```

13.17003734, 0.34421834, 7.27890893, 15.51511559, 0.92800099,
3.63485054, 6.80184256, 11.451762, 10.69664795, 8.59383781,
7.50446583, 10.33994154, 11.85072027, 13.544946, 10.27236403,
9.18410191, 7.7208794, 12.33421272, 6.527156, 11.17483778,
7.97617745, 9.31452692, 9.56473016, 9.51077399, 12.20917888,
12.10672271, 4.70427674, 6.38943267, 10.02410014, 11.97786002,
12.77246335, 6.50139525, 10.64829499, 7.7058727, 6.05475715,
11.28248424, 10.75341994, 17.22835762, 9.53819376, 8.96368426,
6.61412036, 12.00162611, 5.85400348, 4.07058709, 10.08426584,
10.15760235, 11.50892785, 10.58412873, 10.32113545, 12.98841501,
0.40061030, 11.00000000, 6.70010000, 10.00000000, 10.00000000

```

```

In [67]: #prediction in the train data
pred=MLR.predict(x_train)
pred

```

```

Out[67]: array([10.64104453, 11.72955404, 9.71670847, ..., 9.33031288,
11.94411399, 9.8609076 ])

```

```

In [68]: from sklearn.metrics import r2_score

acc=r2_score(y_test,y_pred)

acc

```

```

Out[68]: 0.4331576346139585

```

```

In [69]: #test this model

```

```

Out[68]: 0.4331576346139585

```

```

In [69]: #test this model

MLR.predict([[1,0.455,0.365,0.095,0.5140,0.2245,0.1010,0.150]])

```

```

Out[69]: array([9.91033204])

```

14. Measure the performance using Metrics.¶

```

In [70]: from sklearn import metrics
from sklearn.metrics import mean_squared_error

```

```

In [71]: np.sqrt(mean_squared_error(y_test,y_pred))

```

```

Out[71]: 2.4905110779015462

```

LASSO

```

In [72]: from sklearn.linear_model import Lasso, Ridge

```

```

In [73]: #intialising model

```

```

In [73]: #intialising model

lso=Lasso(alpha=0.01,normalize=True)

```

```

In [74]: #fit the model
lso.fit(x_train,y_train)

```

```

Out[74]: Lasso(alpha=0.01, normalize=True)

```

```

In [75]: #prediction on test data

lso_pred=lso.predict(x_test)

```

```

In [76]: #coef
coef=lso.coef_
coef

```

```

Out[76]: array([-0.         , 0.         , 0.         , 0.4751529 , 0.18634695,
0.         , 0.         , 0.8021721 ])

```

```

In [77]: #accuracy

from sklearn import metrics
from sklearn.metrics import mean_squared_error
metrics.r2_score(y_test,lso_pred)

```

```

Out[77]: 0.3760900261255968

```

Out[78]: 2.715552909824135

```
In [79]: rg=Ridge(alpha=0.01,normalize=True)
```

```
Out[80]: Ridge(alpha=0.01, normalize=True)
```

```
Out[81]: array([11.49838542,  9.22452452,  6.72241086,  7.80010402, 12.09475499,
 11.37301357, 11.13313 ,  8.85299136, 10.95426872,  6.83332623,
 10.48221326,  9.08348674, 12.3998871 , 14.39846005, 13.62343834,
 10.11925891, 11.53997639,  7.75730522, 12.85320604,  8.43018605,
 6.53855123,  8.20224034,  7.58755052, 10.2671289 , 10.65653767,
 11.30141111,  7.50735436,  9.91086293, 11.27856902, 11.29021902,
 10.93344581, 10.32246436, 10.4456454 , 10.47230589, 11.05682097,
 10.1640513 , 10.10050704,  6.5623351 , 11.84100809,  6.75171646,
 4.18665064, 11.0291328 ,  7.72116038,  9.60080953, 11.57691909,
 13.01362452,  6.35434964,  7.30414243, 15.1541625 ,  6.91515291,
 4.16356146,  6.81943931, 11.43766939, 10.62078881,  8.65255458,
 7.53582353, 10.44494347, 11.86697333, 13.45239251, 10.40153892,
 9.1961334 ,  7.75332002, 12.25958727,  6.54710958, 11.17149665,
 7.96864693,  9.37526527,  9.69298327,  9.54666379, 12.19633696,
 12.1127204 ,  4.82993146,  6.43644112,  9.93303646, 12.00590353,
 12.76127566,  6.53636246, 10.58092597,  7.71945979,  5.98307484,
 11.30567761, 10.77297947,  6.83531384,  9.62951405,  9.07577171,
 6.65847637, 11.98056215,  5.83715385,  4.18471904, 10.06479866,
 10.18956629, 11.5066688 , 10.63940289, 10.38796727, 12.94599046,
```

```
In [82]: #coef
         rg.coef_

Out[82]: array([-0.30797338, -0.75443399,  0.34843757,  0.94370227,  0.96851431,
                -1.38791368, -0.04943813,  1.70772786])
```

Out[83]: 0.43177328549243543

Out[84]: 2.4935504011542577

```
In [ ]: # 1. Download the dataset

        #importing the libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
```

```

import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

# 2. Load the dataset into the tool.

#Loading the dataset

d = pd.read_csv(r'Downloads/abalone.csv')

# 3. Perform Below Visualizations.

# • Univariate Analysis

d.head()

#Boxplot

sns.boxplot(d['Diameter'])

#histogram

plt.hist(d['Diameter'])

#Line plot

```

```

#Line plot

plt.plot(d['Diameter'].head(10))

#piechart

plt.pie(d['Diameter'].head(),autopct='%.2f')

#distplot

sns.distplot(d['Diameter'].head(200))

# • Bi - Variate Analysis

#scatter plot

plt.scatter(d['Diameter'].head(500),d['Length'].head(500))

#bar plot

plt.bar(d['Sex'].head(10),d['Rings'].head(10))

#Labelling of x,y and result

plt.title('Bar plot')

```

```

plt.xlabel('Diameter')
plt.ylabel('Rings')

sns.barplot(d['Sex'], d['Rings'])

#joint plot

sns.jointplot(d['Diameter'].head(50),d['Rings'].head(50))

#bar plot

sns.barplot('Diameter','Rings',hue='Sex',data=d.head())

sns.lineplot(d['Diameter'].head(),d['Rings'].head())

# • Multi - Variate Analysis

#boxplot

sns.boxplot(d['Sex'].head(10),d['Diameter'].head(10),d['Rings'].head(10))

#heat map

fig=plt.figure(figsize=(8,5))
sns.heatmap(d.head().corr(),annot=True)

#pair plot

sns.pairplot(d.head(),hue='Rings')

```

```

sns.pairplot(d.head(),hue='Rings')

sns.pairplot(d.head())

# 4. Perform descriptive statistics on the dataset.

#head
d.head()

#tail
d.tail()

d.info()

d.describe()

#mode
d.mode().T

d.shape

#kurtosis

```

```

d.kurt()

#skewness
d.skew()

#variance
d.var()

#finding unique values for columns
d.nunique()

# 5. Check for Missing values and deal with them.

#finding missing values
d.isna()

d.isna().any()

d.isna().sum()

d.isna().any().sum()
#no missing values

```

```

# 6. Find the outliers and replace them outliers

#finding outliers
sns.boxplot(d['Diameter'])

#handling outliers
qnt=d.quantile(q=[0.25,0.75])
qnt

iqr=qnt.loc[0.75]-qnt.loc[0.25]

iqr

lower=qnt.loc[0.25]-(1.5*iqr)
lower

upper=qnt.loc[0.75]+(1.5*iqr)
upper

# replacing outliers

##Diameter
d['Diameter']=np.where(d['Diameter']<0.155,0.4078,d['Diameter'])
sns.boxplot(d['Diameter'])

```

```

## Length
sns.boxplot(d['Length'])

d['Length']=np.where(d['Length']<0.23,0.52, d['Length'])

sns.boxplot(d['Length'])

## Height
sns.boxplot(d['Height'])

d['Height']=np.where(d['Height']<0.04,0.139, d['Height'])
d['Height']=np.where(d['Height']>0.23,0.139, d['Height'])

sns.boxplot(d['Height'])

## Whole weight
sns.boxplot(d['Whole weight'])

d['Whole weight']=np.where(d['Whole weight']>0.9,0.82, d['Whole weight'])

sns.boxplot(d['Whole weight'])

## Shucked weight

```

```

sns.boxplot(d['Shucked weight'])

d['Shucked weight']=np.where(d['Shucked weight']>0.93,0.35, d['Shucked weight'])

sns.boxplot(d['Shucked weight'])

## Viscera weight
sns.boxplot(d['Viscera weight'])

d['Viscera weight']=np.where(d['Viscera weight']>0.46,0.18, d['Viscera weight'])

sns.boxplot(d['Viscera weight'])

## Shell weight
sns.boxplot(d['Shell weight'])

d['Shell weight']=np.where(d['Shell weight']>0.61,0.2388, d['Shell weight'])

sns.boxplot(d['Shell weight'])

# 7. Check for Categorical columns and perform encoding.

#one hot encoding

```

```

# 8. Split the data into dependent and independent variables.

x=d.drop(columns= ['Rings'])
y=d['Rings']
x

y

# 9. Scale the independent variables

from sklearn.preprocessing import scale #StandardScaler

#Scaling the independent variables

x = scale(x)
x

# 10. Split the data into training and testing

from sklearn.model_selection import train_test_split

```

```

# 10. Split the data into training and testing

from sklearn.model_selection import train_test_split

#splitting data to train and test
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2)
print(x_train.shape, x_test.shape)

# 11. Build the Model

#Multiple Regression

from sklearn.linear_model import LinearRegression

MLR=LinearRegression()

# 12. Train the Model

MLR.fit(x_train,y_train)

# 13. Test the Model

#prediction on the test data
y_pred=MLR.predict(x_test)
y_pred

```

```

#prediction in the train data
pred=MLR.predict(x_train)
pred

from sklearn.metrics import r2_score

acc=r2_score(y_test,y_pred)

acc

#test this model

MLR.predict([[1,0.455,0.365,0.095,0.5140,0.2245,0.1010,0.150]])

# 14. Measure the performance using Metrics.¶

from sklearn import metrics
from sklearn.metrics import mean_squared_error

np.sqrt(mean_squared_error(y_test,y_pred))

# LASSO

from sklearn.linear_model import Lasso, Ridge

#initialising model

lso=Lasso(alpha=0.01,normalize=True)

```

```

# LASSO

from sklearn.linear_model import Lasso, Ridge

#initialising model

lso=Lasso(alpha=0.01,normalize=True)

#fit the model
lso.fit(x_train,y_train)

#prediction on test data

lso_pred=lso.predict(x_test)

#coef
coef=lso.coef_
coef

#accuracy

from sklearn import metrics
from sklearn.metrics import mean_squared_error
metrics.r2_score(y_test,lso_pred)

#error

```