

IMPORTING LIBRARIES

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import datetime
from pylab import rcParams
import matplotlib.pyplot as plt
import warnings
import itertools
import statsmodels.api as sm
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers import Dropout
from sklearn.metrics import mean_squared_error
from keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
import seaborn as sns
sns.set_context("paper", font_scale=1.3)
sns.set_style('white')
import math
from sklearn.preprocessing import MinMaxScaler
# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list
warnings.filterwarnings("ignore")
plt.style.use('fivethirtyeight')
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

IMPORTING DATA

```
In [10]: import io
df = pd.read_excel('Crude Oil Prices Daily.xlsx')
df.head()
df[:10]
```

Out[10]:

	Date	Closing Value
--	------	---------------

0	1986-01-02	25.56
---	------------	-------

1	1986-01-03	26.00
---	------------	-------

2	1986-01-06	26.53
---	------------	-------

3	1986-01-07	25.85
---	------------	-------

4	1986-01-08	25.87
---	------------	-------

5	1986-01-09	26.03
---	------------	-------

6	1986-01-10	25.65
---	------------	-------

7	1986-01-13	25.08
---	------------	-------

8	1986-01-14	24.97
---	------------	-------

```
In [11]: #Sort dataset by column Date
df = df.sort_values('Date')
df = df.groupby('Date')['Closing Value'].sum().reset_index()
df.set_index('Date', inplace=True)
df=df.loc[datetime.date(year=2000,month=1,day=1):]
```

```
In [12]: df.head()
```

```
Out[12]:
```

	Closing Value
Date	
2000-01-04	25.56
2000-01-05	24.65
2000-01-06	24.79
2000-01-07	24.79
2000-01-10	24.71

DATA PRE-PROCESSING

```
In [13]: def DfInfo(df_initial):
# gives some infos on columns types and number of null values
tab_info = pd.DataFrame(df_initial.dtypes).T.rename(index={0: 'column type'})
tab_info = tab_info.append(pd.DataFrame(df_initial.isnull().sum()).T.rename(index={0: 'null values (nb)'}))
tab_info = tab_info.append(pd.DataFrame(df_initial.isnull().sum() / df_initial.shape[0]).T.rename(index={0: 'null values (%)'}))

return tab_info
```

```
In [14]: DfInfo(df)
```

```
Out[14]:
```

	Closing Value
column type	float64
null values (nb)	0
null values (%)	0.0


```
In [15]: df.index
```

```
Out[15]: DatetimeIndex(['2000-01-04', '2000-01-05', '2000-01-06', '2000-01-07',
                        '2000-01-10', '2000-01-11', '2000-01-12', '2000-01-13',
                        '2000-01-14', '2000-01-18',
                        ...,
                        '2018-06-26', '2018-06-27', '2018-06-28', '2018-06-29',
                        '2018-07-02', '2018-07-03', '2018-07-04', '2018-07-05',
```

```
'2018-07-06', '2018-07-09'],  
dtype='datetime64[ns]', name='Date', length=4673, freq=None)
```

```
In [16]: y = df['Closing Value'].resample('MS').mean()
```

```
In [17]: y.plot(figsize=(15, 6))  
plt.show()
```



```
In [18]: rcParams['figure.figsize'] = 18, 8  
decomposition = sm.tsa.seasonal_decompose(y, model='additive')  
fig = decomposition.plot()  
plt.show()
```



```
In [19]: sc = MinMaxScaler(feature_range = (0, 1))  
df = sc.fit_transform(df)
```

TRAINING AND TESTING

```
In [20]: train_size = int(len(df) * 0.70)  
test_size = len(df) - train_size  
train, test = df[0:train_size, :], df[train_size:len(df), :]
```

```
In [21]: def create_data_set(_data_set, _look_back=1):
data_x, data_y = [], []
for i in range(len(_data_set) - _look_back - 1):
a = _data_set[i:(i + _look_back), 0]
data_x.append(a)
data_y.append(_data_set[i + _look_back, 0])
return np.array(data_x), np.array(data_y)
```

```
In [23]: look_back = 90
X_train, Y_train, X_test, Y_test = [], [], [], []
X_train, Y_train = create_data_set(train, look_back)
X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
X_test, Y_test = create_data_set(test, look_back)
X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

LSTM LAYER

```
In [24]: regressor = Sequential()

regressor.add(LSTM(units = 60, return_sequences = True, input_shape = (X_train.
regressor.add(Dropout(0.1))

regressor.add(LSTM(units = 60, return_sequences = True))
regressor.add(Dropout(0.1))

regressor.add(LSTM(units = 60))
regressor.add(Dropout(0.1))

regressor.add(Dense(units = 1))

regressor.compile(optimizer = 'adam', loss = 'mean_squared_error')
reduce_lr = ReduceLROnPlateau(monitor='val_loss', patience=5)
history = regressor.fit(X_train, Y_train, epochs = 20, batch_size = 15, validation
```

```
Epoch 1/20
212/212 [=====] - 22s 85ms/step - loss: 0.0045 - val_loss: 0.0321 - lr: 0.0010
Epoch 2/20
212/212 [=====] - 17s 79ms/step - loss: 0.0122 - val_loss: 0.0445 - lr: 0.0010
Epoch 3/20
212/212 [=====] - 17s 81ms/step - loss: 0.0110 - val_loss: 0.0513 - lr: 0.0010
Epoch 4/20
212/212 [=====] - 17s 78ms/step - loss: 0.0164 - val_loss: 0.0483 - lr: 0.0010
Epoch 5/20
212/212 [=====] - 17s 79ms/step - loss: 0.0187 - val_loss: 0.0549 - lr: 0.0010
Epoch 6/20
212/212 [=====] - 17s 79ms/step - loss: 0.0185 - val_loss: 0.0478 - lr: 0.0010
Epoch 7/20
212/212 [=====] - 17s 80ms/step - loss: 0.0186 - val_loss: 0.0031 - lr: 1.0000e-04
Epoch 8/20
```



```

212/212 [=====] - 17s 80ms/step - loss: 0.0032 - val_l
oss: 0.0024 - lr: 1.0000e-04
Epoch 9/20
212/212 [=====] - 17s 79ms/step - loss: 0.0024 - val_l
oss: 0.0019 - lr: 1.0000e-04
Epoch 10/20
212/212 [=====] - 17s 79ms/step - loss: 0.0021 - val_l
oss: 0.0017 - lr: 1.0000e-04
Epoch 11/20
212/212 [=====] - 17s 82ms/step - loss: 0.0016 - val_l
oss: 0.0017 - lr: 1.0000e-04
Epoch 12/20
212/212 [=====] - 17s 79ms/step - loss: 0.0015 - val_l
oss: 0.0016 - lr: 1.0000e-04
Epoch 13/20
212/212 [=====] - 17s 80ms/step - loss: 0.0013 - val_l
oss: 0.0014 - lr: 1.0000e-04
Epoch 14/20
212/212 [=====] - 17s 79ms/step - loss: 0.0012 - val_l
oss: 0.0013 - lr: 1.0000e-04
Epoch 15/20
212/212 [=====] - 17s 78ms/step - loss: 0.0011 - val_l
oss: 0.0013 - lr: 1.0000e-04
Epoch 16/20
212/212 [=====] - 17s 78ms/step - loss: 9.6755e-04 - v
al_loss: 0.0013 - lr: 1.0000e-04
Epoch 17/20
212/212 [=====] - 17s 80ms/step - loss: 0.0010 - val_l
oss: 0.0013 - lr: 1.0000e-04
Epoch 18/20
212/212 [=====] - 17s 79ms/step - loss: 9.4863e-04 - v
al_loss: 0.0013 - lr: 1.0000e-04
Epoch 19/20
212/212 [=====] - 17s 81ms/step - loss: 9.1643e-04 - v
al_loss: 0.0014 - lr: 1.0000e-04
Epoch 20/20
212/212 [=====] - 18s 84ms/step - loss: 9.3516e-04 - v
al_loss: 0.0013 - lr: 1.0000e-04

```

MODEL TRAINING

```

In [25]: train_predict = regressor.predict(X_train)
         test_predict = regressor.predict(X_test)

```

```

100/100 [=====] - 4s 27ms/step
41/41 [=====] - 1s 28ms/step

```

```

In [26]: train_predict = sc.inverse_transform(train_predict)
         Y_train = sc.inverse_transform([Y_train])
         test_predict = sc.inverse_transform(test_predict)
         Y_test = sc.inverse_transform([Y_test])

```

PREDICTION

```

In [27]: print('Train Mean Absolute Error:', mean_absolute_error(Y_train[0], train_predi
         print('Train Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_train[0], t
         print('Test Mean Absolute Error:', mean_absolute_error(Y_test[0], test_predict[

```

```

print('Test Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_test[0], tes
plt.figure(figsize=(8,4))
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epochs')
plt.legend(loc='upper right')
plt.show();

```

Train Mean Absolute Error: 2.7544575163643317

Train Root Mean Squared Error: 3.5074279586248873

Test Mean Absolute Error: 2.373878536093426

Test Root Mean Squared Error: 5.286033010892939



In [28]:

```

aa=[x for x in range(180)]
plt.figure(figsize=(8,4))
plt.plot(aa, Y_test[0][:180], marker='.', label="actual")
plt.plot(aa, test_predict[:,0][:180], 'r', label="prediction")
plt.tight_layout()
sns.despine(top=True)
plt.subplots_adjust(left=0.07)
plt.ylabel('Price', size=15)
plt.xlabel('Time step', size=15)
plt.legend(fontsize=15)
plt.show();

```

